

# LogisticR

November 21, 2024

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: df = pd.read_csv('heart.csv')

# Print the number of records with and without heart disease.
print("Number of records in each label are")
print(df['target'].value_counts())

# Print the percentage of each label
print("\nPercentage of records in each label are")
print(df['target'].value_counts() * 100 / df.shape[0], "\n")

# Print the first five rows of Dataframe.
df.head()
```

Number of records in each label are

target

1 165

0 138

Name: count, dtype: int64

Percentage of records in each label are

target

1 54.455446

0 45.544554

Name: count, dtype: float64

```
[3]: Unnamed: 0  age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  \
0           0   63   1   3        145   233   1         0        150     0
1           1   37   1   2        130   250   0         1        187     0
2           2   41   0   1        130   204   0         0        172     0
3           3   56   1   1        120   236   0         1        178     0
4           4   57   0   0        120   354   0         1        163     1
```


	oldpeak	slope	ca	thal	target
0	2.3	0	0	1	1
1	3.5	0	0	2	1
2	1.4	2	0	2	1
3	0.8	2	0	2	1
4	0.6	2	0	2	1

```
[4]: def sigmoid(x):
      return pd.Series(1 / ( 1 + np.exp(-x)))
```

```
[5]: df['chol'].describe()
```

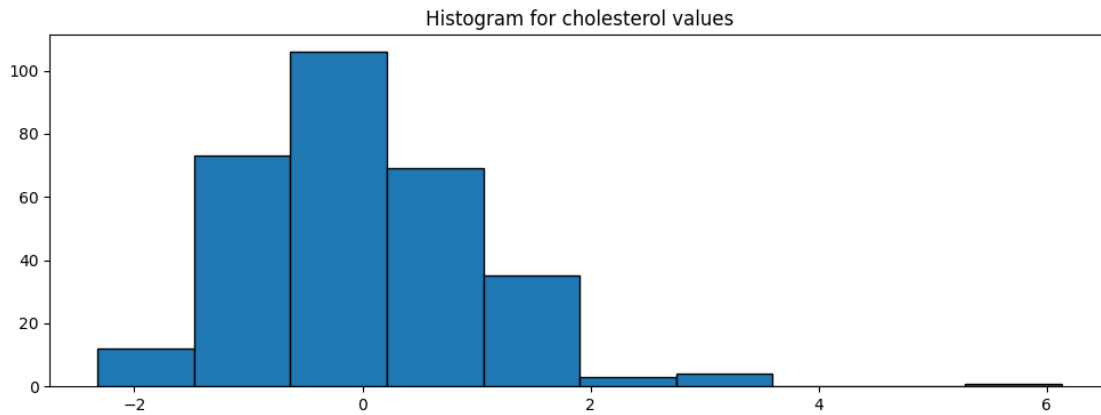
```
[5]: count    303.000000
      mean     246.264026
      std      51.830751
      min     126.000000
      25%     211.000000
      50%     240.000000
      75%     274.500000
      max     564.000000
      Name: chol, dtype: float64
```

```
[6]: plt.figure(figsize = (12,4), dpi = 96)
      plt.title("Histogram for cholesterol values")
      plt.hist(df['chol'], bins = 'sturges', edgecolor = 'black')
      plt.show()
```



```
[7]: def standard_scalar(series):
      new_series = (series - series.mean()) / series.std()
      return new_series
      scaled_chol = standard_scalar(df['chol'])
```

```
plt.figure(figsize = (12,4))
plt.title("Histogram for cholesterol values")
plt.hist(scaled_chol, bins = 'sturges', edgecolor = 'black')
plt.show()
```



```
[8]: chol_sig_output = sigmoid(df['chol'])
chol_sig_output.describe()
```

```
[8]: count    303.0
mean         1.0
std          0.0
min          1.0
25%          1.0
50%          1.0
75%          1.0
max          1.0
Name: chol, dtype: float64
```


```
[9]: scaled_chol_sig_output = sigmoid(scaled_chol)
scaled_chol_sig_output.describe()
```

```
[9]: count    303.000000
mean     0.492837
std     0.198175
min     0.089454
25%     0.336179
50%     0.469823
75%     0.632919
max     0.997829
Name: chol, dtype: float64
```

```
[10]: def predict(sig_output, threshold):
        y_pred = [ 1 if output >= threshold else 0 for output in sig_output]
        return pd.Series(y_pred)
```

```
[11]: threshold = 0.5
heart_disease_pred = predict(scaled_chol_sig_output, threshold)

plt.figure(figsize=(13,3), dpi = 96)
plt.scatter(scaled_chol, heart_disease_pred)
plt.axhline(y = threshold, label = f'y = { threshold }', color = 'r')
plt.legend()
plt.show()
```



```
[12]: print(f"Threshold value: {threshold}")
print(f"\nPredicted value counts:\n{heart_disease_pred.value_counts()}")
print(f"\nActual value counts:\n{df['target']. value_counts()}")
```

Threshold value: 0.5

Predicted value counts:

0     167

1     136

Name: count, dtype: int64

Actual value counts:

target

1     165

0     138

Name: count, dtype: int64

```
[13]: from sklearn.metrics import confusion_matrix

print(confusion_matrix(df['target'], heart_disease_pred))
```

[[ 65 73]

[102 63]]

```
[14]: from sklearn.metrics import classification_report

print(classification_report(df['target'], heart_disease_pred))
```

	precision	recall	f1-score	support
0	0.39	0.47	0.43	138
1	0.46	0.38	0.42	165
accuracy			0.42	303
macro avg	0.43	0.43	0.42	303
weighted avg	0.43	0.42	0.42	303

```
[15]: #Split the training and testing data
from sklearn.model_selection import train_test_split

X = df.drop(columns = 'target')
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳random_state = 42)
```

```
[16]: from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

# Standardizing the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize Logistic Regression with increased max_iter
log_clf_1 = LogisticRegression(max_iter=200)

# Train the model on the scaled data
log_clf_1.fit(X_train_scaled, y_train)

# Print the training accuracy
print(f"Training Accuracy: {log_clf_1.score(X_train_scaled, y_train):.4f}")

# Predict the target values for the train set
y_train_pred = log_clf_1.predict(X_train_scaled)

# Confusion Matrix
print("\nConfusion Matrix:\n")
print(confusion_matrix(y_train, y_train_pred))

# Classification Report
```

```
print("\nClassification Report:\n")
print(classification_report(y_train, y_train_pred))
```

Training Accuracy: 0.9953

Confusion Matrix:

```
[[ 96   1]
 [  0 115]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	97
1	0.99	1.00	1.00	115
accuracy			1.00	212
macro avg	1.00	0.99	1.00	212
weighted avg	1.00	1.00	1.00	212

```
[17]: #Normalise the train and test data-frames using the standard normalisation
      ↪method.
def standard_scaler(series):
    new_series = (series - series.mean()) / series.std()
    return new_series

norm_X_train = X_train.apply(standard_scaler, axis = 0)
norm_X_test = X_test.apply(standard_scaler, axis = 0)

norm_X_train.describe()
```

```
[17]:
```

	Unnamed: 0	age	sex	cp	trestbps	\
count	2.120000e+02	2.120000e+02	2.120000e+02	2.120000e+02	2.120000e+02	
mean	1.508228e-16	1.864337e-16	1.298751e-16	2.251867e-17	5.697748e-16	
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	
min	-1.728945e+00	-2.757098e+00	-1.391141e+00	-9.778484e-01	-2.142798e+00	
25%	-8.899408e-01	-7.177485e-01	-1.391141e+00	-9.778484e-01	-6.152369e-01	
50%	2.326080e-02	7.080006e-02	7.154438e-01	-1.364440e-02	-2.771338e-02	
75%	8.308735e-01	7.233920e-01	7.154438e-01	9.505596e-01	5.598102e-01	
max	1.718391e+00	2.463637e+00	7.154438e-01	1.914764e+00	3.614933e+00	

	chol	fbs	restecg	thalach	exang	\
count	2.120000e+02	2.120000e+02	2.120000e+02	2.120000e+02	2.120000e+02	
mean	1.424437e-16	-5.812960e-17	-1.005485e-16	3.058350e-16	9.216946e-17	
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	

min	-2.129975e+00	-3.811266e-01	-1.029172e+00	-2.731467e+00	-6.855616e-01
25%	-6.649586e-01	-3.811266e-01	-1.029172e+00	-6.547229e-01	-6.855616e-01
50%	-1.338901e-01	-3.811266e-01	8.680843e-01	1.693821e-01	-6.855616e-01
75%	5.162111e-01	-3.811266e-01	8.680843e-01	7.847138e-01	1.451778e+00
max	5.799427e+00	2.611423e+00	2.765341e+00	2.279091e+00	1.451778e+00

	oldpeak	slope	ca	thal
count	2.120000e+02	2.120000e+02	2.120000e+02	2.120000e+02
mean	7.541138e-17	5.865329e-17	7.960090e-17	3.770569e-17
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-9.289910e-01	-2.305793e+00	-6.746937e-01	-3.912465e+00
25%	-9.289910e-01	-6.763660e-01	-6.746937e-01	-5.475864e-01
50%	-1.961683e-01	-6.763660e-01	-6.746937e-01	-5.475864e-01
75%	5.366543e-01	9.530612e-01	3.770347e-01	1.134853e+00
max	4.200768e+00	9.530612e-01	3.532220e+00	1.134853e+00

```
[18]: norm_X_test.describe()
```

```
[18]: Unnamed: 0      age      sex      cp      trestbps \
count  9.100000e+01  9.100000e+01  9.100000e+01  9.100000e+01  9.100000e+01
mean   -1.249001e-16 -2.147245e-16 -1.390829e-16 -1.952040e-17 -6.868742e-16
std     1.000000e+00  1.000000e+00  1.000000e+00  1.000000e+00  1.000000e+00
min    -1.644808e+00 -2.301763e+00 -1.661622e+00 -8.425578e-01 -1.853721e+00
25%    -8.333169e-01 -8.354271e-01 -1.661622e+00 -8.425578e-01 -6.650121e-01
50%    -6.722410e-02  1.797284e-01  5.952080e-01 -8.425578e-01 -1.662530e-02
75%     9.258592e-01  6.309086e-01  5.952080e-01  1.123410e+00  4.696648e-01
max     1.703301e+00  2.435630e+00  5.952080e-01  2.106394e+00  3.549502e+00
```

	chol	fbs	restecg	thalach	exang
count	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01
mean	-4.148086e-17	3.538073e-17	-4.880101e-18	-5.294910e-16	-1.049222e-16
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-2.624853e+00	-4.938276e-01	-9.430373e-01	-3.319275e+00	-7.148350e-01
25%	-7.201088e-01	-4.938276e-01	-9.430373e-01	-6.418709e-01	-7.148350e-01
50%	-1.836075e-02	-4.938276e-01	-9.430373e-01	1.078023e-01	-7.148350e-01
75%	6.165541e-01	-4.938276e-01	9.639937e-01	6.432832e-01	1.383552e+00
max	3.679740e+00	2.002745e+00	2.871025e+00	1.864180e+00	1.383552e+00

	oldpeak	slope	ca	thal
count	9.100000e+01	9.100000e+01	9.100000e+01	9.100000e+01
mean	1.339037e-16	-1.244426e-16	-6.344132e-17	1.848338e-16
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-8.367971e-01	-2.184053e+00	-8.102615e-01	-3.491486e+00
25%	-8.367971e-01	-5.812398e-01	-8.102615e-01	-4.364358e-01
50%	-3.799059e-01	-5.812398e-01	-8.102615e-01	-4.364358e-01
75%	5.719508e-01	1.021573e+00	9.246514e-01	1.091089e+00
max	3.884412e+00	1.021573e+00	2.659564e+00	1.091089e+00

```
[19]: #Create a dictionary containing the different combination of features selected
      ↪by RFE and their corresponding f1-scores.
      # Import the libraries
      from sklearn.feature_selection import RFE
      from sklearn.metrics import f1_score
      from sklearn.linear_model import LogisticRegression

      # Create the empty dictionary.
      dict_rfe = {}

      # Create a loop
      for i in range(1, len(X_train.columns) + 1):
          lg_clf_2 = LogisticRegression()
          rfe = RFE(lg_clf_2, n_features_to_select=i) # 'i' is the number of features to
          ↪be selected by RFE to fit a logistic regression model on norm_X_train and
          ↪y_train.
          rfe.fit(norm_X_train, y_train)

          rfe_features = list(norm_X_train.columns[rfe.support_]) # A list of important
          ↪features chosen by RFE.
          rfe_X_train = norm_X_train[rfe_features]

          # Build a logistic regression model using the features selected by RFE.
          lg_clf_3 = LogisticRegression()
          lg_clf_3.fit(rfe_X_train, y_train)

          # Predicting 'y' values only for the test set as generally, they are
          ↪predicted quite accurately for the train set.
          y_test_pred = lg_clf_3.predict(norm_X_test[rfe_features])

          f1_scores_array = f1_score(y_test, y_test_pred, average = None)
          dict_rfe[i] = {"features": list(rfe_features), "f1_score": f1_scores_array} #
          ↪'i' is the number of features to be selected by RFE.
```

```
[20]: #Print the dictionary created
      dict_rfe
```

```
[20]: {1: {'features': ['Unnamed: 0'], 'f1_score': array([0.98795181, 0.98989899])},
      2: {'features': ['Unnamed: 0', 'oldpeak'],
          'f1_score': array([0.98765432, 0.99009901])},
      3: {'features': ['Unnamed: 0', 'exang', 'oldpeak'],
          'f1_score': array([0.98765432, 0.99009901])},
      4: {'features': ['Unnamed: 0', 'exang', 'oldpeak', 'thal'],
          'f1_score': array([0.97560976, 0.98      ])},
      5: {'features': ['Unnamed: 0', 'restecg', 'exang', 'oldpeak', 'thal'],
          'f1_score': array([0.97560976, 0.98      ])},
      6: {'features': ['Unnamed: 0', 'sex', 'restecg', 'exang', 'oldpeak', 'thal'],
```



```

    'f1_score': array([0.96385542, 0.96969697])},
7: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'restecg',
    'exang',
    'oldpeak',
    'thal'],
    'f1_score': array([0.97560976, 0.98      ])},
8: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'restecg',
    'exang',
    'oldpeak',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
9: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'restecg',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
10: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'chol',
    'restecg',
    'exang',
    'oldpeak',
    'slope',
    'ca',
    'thal'],
    'f1_score': array([0.96385542, 0.96969697])},
11: {'features': ['Unnamed: 0',
    'sex',
    'cp',
    'trestbps',
    'chol',
    'restecg',
    'exang',
    'oldpeak',

```

```

'slope',
'ca',
'thal'],
'f1_score': array([0.96385542, 0.96969697])},
12: {'features': ['Unnamed: 0',
'sex',
'cp',
'trestbps',
'chol',
'fbs',
'restecg',
'exang',
'oldpeak',
'slope',
'ca',
'thal'],
'f1_score': array([0.96385542, 0.96969697])},
13: {'features': ['Unnamed: 0',
'sex',
'cp',
'trestbps',
'chol',
'fbs',
'restecg',
'thalach',
'exang',
'oldpeak',
'slope',
'ca',
'thal'],
'f1_score': array([0.96385542, 0.96969697])},
14: {'features': ['Unnamed: 0',
'age',
'sex',
'cp',
'trestbps',
'chol',
'fbs',
'restecg',
'thalach',
'exang',
'oldpeak',
'slope',
'ca',
'thal'],
'f1_score': array([0.96385542, 0.96969697])}}

```

```
[21]: #Convert the dictionary to the dataframe
pd.options.display.max_colwidth = 100
f1_df = pd.DataFrame.from_dict(dict_rfe, orient = 'index')
f1_df
```

```
[21]:          features \
1
[Unnamed: 0]
2
[Unnamed: 0, oldpeak]
3
[Unnamed: 0, exang, oldpeak]
4
[Unnamed: 0, exang, oldpeak, thal]
5
[Unnamed: 0, restecg, exang, oldpeak, thal]
6
[Unnamed: 0, sex, restecg, exang, oldpeak, thal]
7
[Unnamed: 0, sex, cp, restecg, exang, oldpeak, thal]
8
[Unnamed: 0, sex, cp, restecg, exang, oldpeak, ca, thal]
9
[Unnamed: 0, sex, cp, restecg, exang, oldpeak, slope, ca, thal]
10
[Unnamed: 0, sex, cp, chol, restecg, exang, oldpeak, slope, ca, thal]
11
[Unnamed: 0, sex, cp, trestbps, chol, restecg, exang, oldpeak, slope, ca, thal]
12
[Unnamed: 0, sex, cp, trestbps, chol, fbs, restecg, exang, oldpeak, slope, ca, thal]
13
[Unnamed: 0, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal]
14
[Unnamed: 0, age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal]

          f1_score
1  [0.9879518072289156, 0.98989898989899]
2  [0.9876543209876543, 0.9900990099009901]
3  [0.9876543209876543, 0.9900990099009901]
4  [0.975609756097561, 0.98]
5  [0.975609756097561, 0.98]
6  [0.963855421686747, 0.9696969696969697]
7  [0.975609756097561, 0.98]
8  [0.963855421686747, 0.9696969696969697]
9  [0.963855421686747, 0.9696969696969697]
10 [0.963855421686747, 0.9696969696969697]
11 [0.963855421686747, 0.9696969696969697]
```

```
12 [0.963855421686747, 0.9696969696969697]
13 [0.963855421686747, 0.9696969696969697]
14 [0.963855421686747, 0.9696969696969697]
```

```
[22]: #Logistic Regression with the ideal number of features.
lg_clf_4 = LogisticRegression()
rfe = RFE(lg_clf_4, n_features_to_select = 3)

rfe.fit(norm_X_train, y_train)

rfe_features = norm_X_train.columns[rfe.support_]
print(rfe_features)
final_X_train = norm_X_train[rfe_features]

lg_clf_4 = LogisticRegression()
lg_clf_4.fit(final_X_train, y_train)

y_test_predict = lg_clf_4.predict(norm_X_test[rfe_features])
final_f1_scores_array = f1_score(y_test, y_test_predict, average = None)
print(final_f1_scores_array)
```

```
Index(['Unnamed: 0', 'exang', 'oldpeak'], dtype='object')
[0.98765432 0.99009901]
```